

PATTERN RECOGNITION AND MACHINE LEARNING

Slide Set 7: Random forests and ensemble methods

November 2019

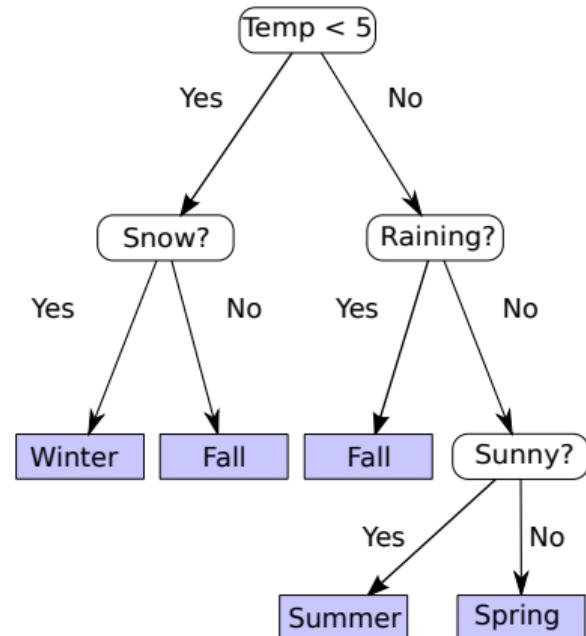
Heikki Huttunen
heikki.huttunen@tuni.fi
Signal Processing
Tampere University

RANDOM FOREST

(and other ensemble methods)

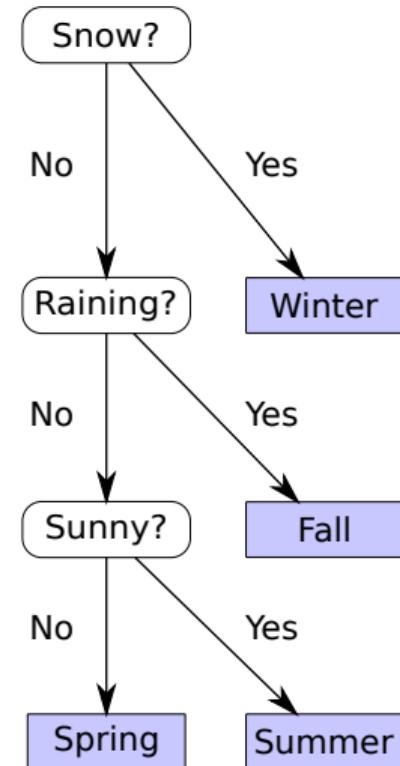
Random Forest

- Random forest (RF) is popular tree based classifier, proposed by Breiman in 1993.
- The RF is based on a collection of *decision trees*.
- A decision tree is a straightforward if-then-else-like diagram that combines individual attributes into a decision.
- Decision trees are easily trained to learn the data.
- For example, the tree on the right predicts the time of year from the following measured attributes: {Temperature, Snow, Rain, Sunny}.



Random Forest

- The problem of decision trees is that they *overlearn* the data, *i.e.*, the training data is memorized with a poor ability to generalize.
- With no restrictions, the DT will have one root-leaf-path per sample, which is essentially same as nearest neighbor.
- RF avoids this by training many "imperfect" trees.
- Each tree is trained with a subset of samples and a subset of features, *i.e.*, some of the attributes are hidden from the training.



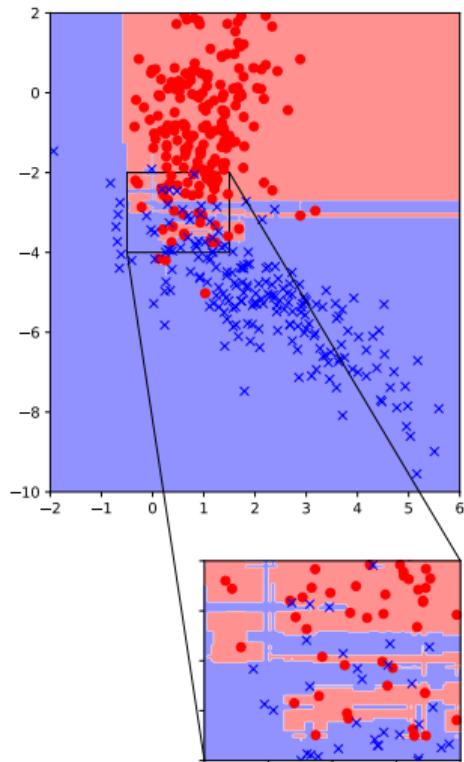
Decision Tree without the *Temperature* attribute.

Random Forest

- The RF extracts values at randomly selected coordinates.
- For example, if we have a data matrix $\mathbf{X} \in \mathbb{R}^{10 \times 4}$ (*i.e.*, 10 samples with 4 features each), we might train trees with the following subsets of the data:
 - **Tree 1:** Train using rows {6,4,7,2,6,3} and columns {1,2,4}
 - **Tree 2:** Train using rows {1,10,2,1,7,9} and columns {1,3}
 - **Tree 3:** Train using rows {7,2,7,3,9,3} and columns {2}
 - ...
- Note that rows (samples) are sampled *with replacement*, *i.e.*, rows may appear more than once.
- Columns are sampled *without replacement* (it does not make sense to repeat the same data).

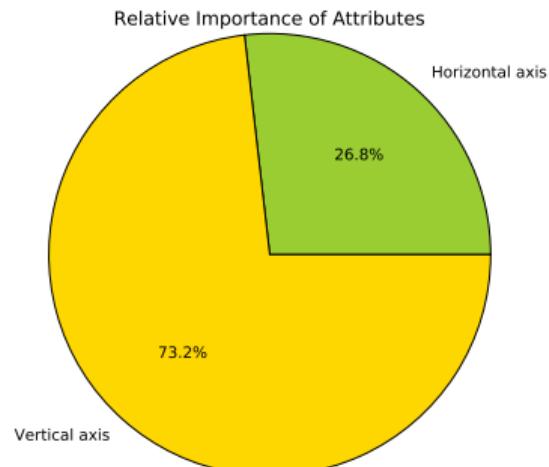
Random Forest

- After training many trees each with different samples and features, the RF predicts the class by taking the majority vote: what is the most frequent label.
- The number of trees varies from a few dozen to few thousand—default in Python is 10.
- The attached picture is produced by a 10-tree random forest.

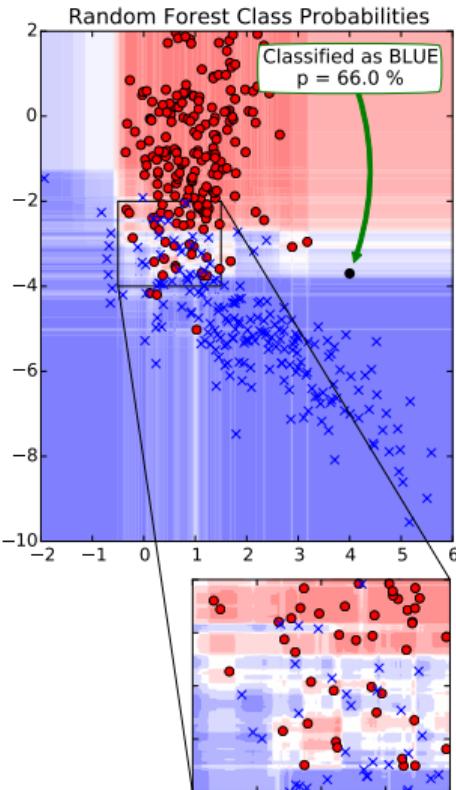


Random Forest

- The collection of trees gives a natural way of estimating class probabilities: Just use the proportion of trees voting for each class.
- RF's also includes a method for assessing feature importances: randomly shuffle each feature at a time and test how much the accuracy drops.



- Loosing an important feature drops the accuracy a lot.
- Shuffling a non-important feature will not change the result much.



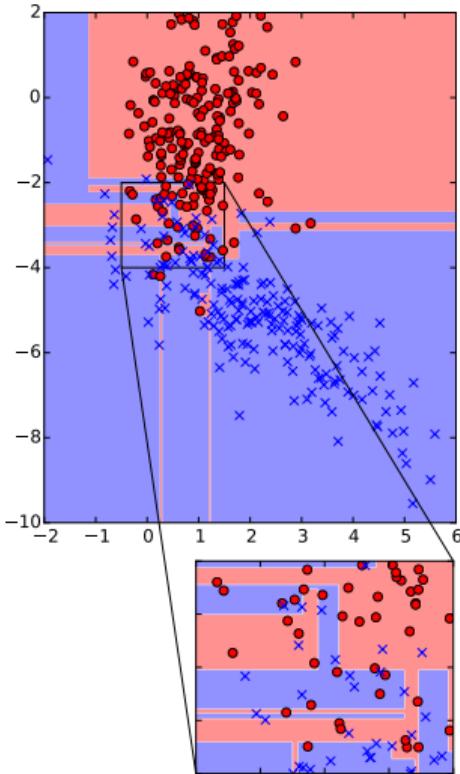
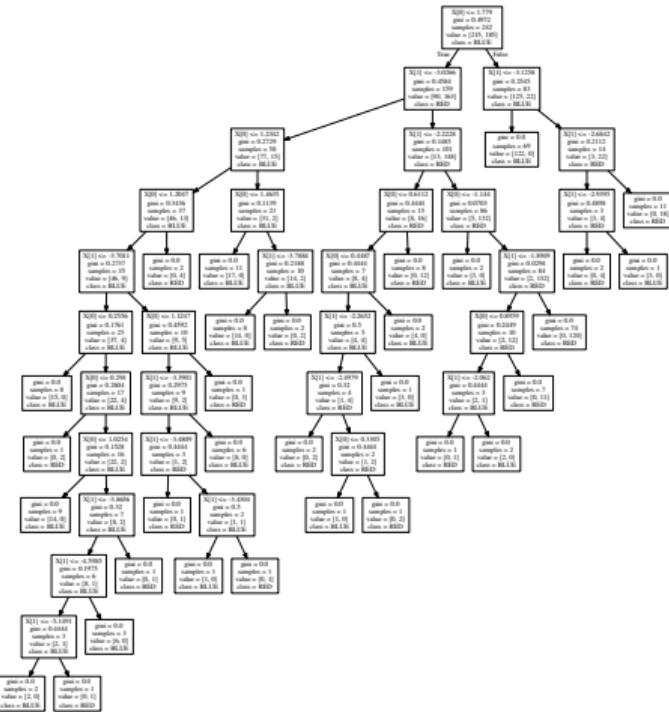
Random Forest in Scikit-Learn

```
# Training code:  
from sklearn.ensemble import RandomForestClassifier  
  
clf = RandomForestClassifier()  
clf.fit(X, y)
```

```
# Testing code:  
>>> clf.predict([0, -4])  
array([ 0.])  
  
>>> clf.predict([0, -2])  
array([ 1.])  
  
>>> clf.predict_proba([[0, -4], [0, -2]])  
array([[ 0.9,  0.1],  
       [ 0.4,  0.6]])  
  
>>> len(clf.estimators_)  
10  
  
>>> type(clf.estimators_[0])  
sklearn.tree.tree.DecisionTreeClassifier
```

- The RandomForestClassifier class is used via the normal interface (.fit() and .predict()).
- Individual trees can be accessed, as well: clf.estimators_ is a list of DecisionTreeClassifier objects.
- One of the 10 trained trees is visualized on the next slide (plot created using sklearn.tree.export_graphviz).

Decision Tree



Other Ensemble Classifiers

- Since the introduction of Random Forest by Breiman in 1993, several extensions have been proposed.
- As a group, these are called *ensemble methods*, because they all consist of an ensemble of weak classifiers,
- Most important ones are briefly summarized in the following slides.

AdaBoost Paradigm

- The **AdaBoost** paradigm was proposed by Freund and Schapire in 1995.
- Also in this case, the classifier consists of a collection of weak classifiers (most often decision trees).
- The difference to other ensemble methods is that trees are grown sequentially:
 - ① Assign each sample a weight $w_1 \dots, w_N$.
 - ② Grow a tree minimizing the classification error weighted by w_n . That is, we emphasize the samples with large w_n .
 - ③ Append the new tree to our ensemble \mathcal{E} .
 - ④ Increase the weights for those samples that were incorrectly classified by \mathcal{E} .
 - ⑤ If not enough trees, then return to step 2.
- Implemented as `sklearn.ensemble.AdaBoostClassifier`.

Gradient Boosted Regression Trees

- **Gradient Boosted Regression Trees** were proposed by Friedman in 1999.
- Another boosting algorithm similar to AdaBoost.
- In this case, the training sequence is the following.
 - 1 Initialize the ensemble \mathcal{E} by a single decision tree fit to the data.
 - 2 Train another tree for correcting the errors made by \mathcal{E} , i.e., attempt to predict $\mathbf{y} - F_{\mathcal{E}}(\mathbf{X})$.
 - 3 Append the new tree to our ensemble \mathcal{E} .
 - 4 If not enough trees, then return to step 2.
- Implemented as `sklearn.ensemble.GradientBoostingClassifier`.

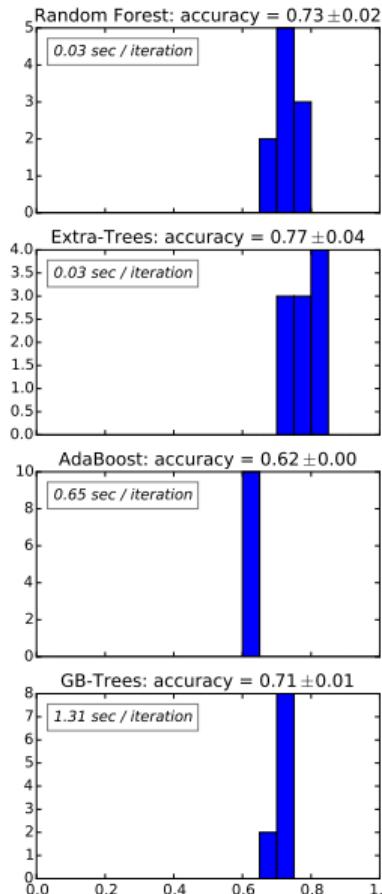
Extremely Randomized Trees

- **Extremely Randomized Trees** were proposed by Geurts *et al.* in 2006.
- The idea is to make the trees even weaker, but to compensate that with the large number of estimators in the ensemble.
- More specifically, not only the features and samples shown to the tree are randomized, but also the growing of individual trees is random.
- In particular the *split point* i.e., the thresholds of comparisons like $X[1] \leq -2.7025$ is randomized (see the graph of a decision tree at an earlier slide).
- Implemented as `sklearn.ensemble.ExtraTreesClassifier`.

Comparison of Methods

- The four ensemble methods were compared with the **arcene dataset**: <https://archive.ics.uci.edu/ml/datasets/Arcene>.
- For randomized algorithms, we iterate the experiment 100 times.

```
# Load Arcene data; 100+100 samples with dimension 10000:  
# Mass spectrometer measurements from ovarian cancer patients and healthy controls.  
X_train, y_train, X_test, y_test = load_arcene()  
  
classifiers = [(RandomForestClassifier(), "Random Forest"),  
               (ExtraTreesClassifier(), "Extra-Trees"),  
               (AdaBoostClassifier(), "AdaBoost"),  
               (GradientBoostingClassifier(), "GB-Trees")]  
  
for clf, name in classifiers:  
    clf.n_estimators = 100  
  
    accuracies = []  
    for iteration in range(100):  
        clf.fit(X_train, y_train)  
        y_hat = clf.predict(X_test)  
        accuracy = accuracy_score(y_test, y_hat)  
        accuracies.append(accuracy)
```



Comparison of Methods

- Let's add a bunch of other classifiers into our for loop.
 - *1-Nearest Neighbor*: **0.88**
 - *Logistic Regression*: 0.84
 - *Linear SVM*: 0.83
 - *5-Nearest Neighbor*: 0.82
 - *LDA*: 0.79
 - *Extra-Trees*: 0.77 ± 0.04
 - *9-Nearest Neighbor*: 0.73
 - *Random Forest*: 0.73 ± 0.02
 - *GB-Trees*: 0.71 ± 0.01
 - *AdaBoost*: 0.62
 - *SVM with RBF kernel*: 0.56
- It seems that linear models and NN excel with this data.